

Naive Bayes Implementation

Author: Srikar Kalle
Student ID: C00313529

Change Log

SL No.	Change Category	Description	Duration (mins)	Difficulty (1-10)
1	Dataset Change	Changed from 'Adult Dataset' to 'Penguins Dataset'.	10	3
2	Feature Selection	Selected <code>bill_length_mm</code> and <code>bill_depth_mm</code> as features.	8	4
3	Target Encoding	Used <code>LabelEncoder</code> to encode the species column.	5	3
4	Data Preprocessing	Standardized features using <code>StandardScaler</code> .	12	5
5	Train-Test Split	Applied an 80-20 split for training and testing.	7	4
6	Model Implementation	Implemented Naïve Bayes (GaussianNB) for classification.	15	6
7	Model Evaluation	Used accuracy, classification report for evaluation.	10	5

Summary

- Transitioned from Kaggle's Adult dataset to Seaborn's Penguins dataset.
- Implemented a Naïve Bayes classification model.
- Improved data preprocessing with scaling.

Performance Metrics

- Model Accuracy: `0.9254` (Generated during execution)

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from datetime import datetime
```

```
In [2]: # Step 2: Data Understanding
df = sns.load_dataset("penguins").dropna()
```

```
In [3]: # Selecting features and target
selected_features = ["bill_length_mm", "bill_depth_mm"]
```

```
X = df[selected_features]
y = LabelEncoder().fit_transform(df["species"]) # Encoding categorical target
```

```
In [4]: # Step 3: Data Preparation
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=
```

```
In [5]: # Step 4: Modeling
nb_model = GaussianNB()
nb_model.fit(X_train, y_train)
```

```
Out[5]: GaussianNB
GaussianNB()
```

```
In [6]: # Step 5: Evaluation
y_pred = nb_model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.4f}")
print(classification_report(y_test, y_pred))
```

Model Accuracy: 0.9254

	precision	recall	f1-score	support
0	0.97	1.00	0.98	31
1	0.79	0.85	0.81	13
2	0.95	0.87	0.91	23
accuracy			0.93	67
macro avg	0.90	0.91	0.90	67
weighted avg	0.93	0.93	0.93	67